



---

# INFYMA HACKATHON '25

---

Muhammad Zeeshan  
Muhammad Essa Zeeshan  
Manshah Hussain



MARCH 10, 2025  
THE OVERFITTERS  
COMSATS University Islamabad

# Report on CNN Model for Image Classification

## 1. Introduction

This project focuses on developing a convolutional neural network (CNN) for classifying images into 5 categories. The model leverages transfer learning using a pre-trained ResNet152V2 base, combined with custom layers for fine-tuning. The approach takes advantage of data augmentation, dropout, and batch normalization to enhance generalization and performance.

## 2. Data Preparation

### Data Augmentation & Preprocessing

To improve the robustness of the model and mitigate overfitting, the following preprocessing steps were applied using Keras' ImageDataGenerator:

- **Training Data:**
  - **Rescaling:** Pixel values normalized by dividing by 255.
  - **Augmentations:**
    - Rotation up to 30 degrees.
    - Horizontal flipping.
    - Shear transformation (range of 0.3).
    - Width and height shifts (range of 0.2).
    - Zoom (range of 0.1).
    - Filling mode set to 'nearest' to handle empty areas after transformations.
- **Validation & Test Data:**
  - **Rescaling Only:** Images are normalized by dividing pixel values by 255 without additional augmentations.

The images are resized to a target size of (224, 224) and are processed in batches of 64.

### 3. Model Architecture

#### Base Model: ResNet152V2

- **Pre-trained on ImageNet:**

The ResNet152V2 model is imported without its top classification layer, allowing the model to utilize pre-learned feature extraction.

- **Layer Freezing:**

All layers except the last 40 layers are frozen. This strategy retains most of the learned features while fine-tuning the deeper layers to better adapt to the new dataset.

#### Custom Layers

Following the base model, additional layers were added:

1. **Dropout (0.5):** Helps prevent overfitting by randomly setting input units to 0.
2. **Flatten:** Transforms the multidimensional output of the base model into a one-dimensional vector.
3. **Batch Normalization:** Normalizes the outputs, improving training stability.
4. **Dense Layer (256 units, He initialization):**
  - Followed by batch normalization, ReLU activation, and dropout.
5. **Dense Layer (128 units, He initialization):**
  - Similarly followed by batch normalization, ReLU activation, and dropout.
6. **Dense Layer (32 units, He initialization):**
  - Followed by batch normalization and ReLU activation.
7. **Output Layer:**
  - A Dense layer with 5 units and softmax activation to predict class probabilities.

#### Parameter Summary:

- **Total Parameters:** ~84.46 million
- **Trainable Parameters:** ~41.16 million
- **Non-trainable Parameters:** ~43.30 million

## 4. Frontend Implementation:

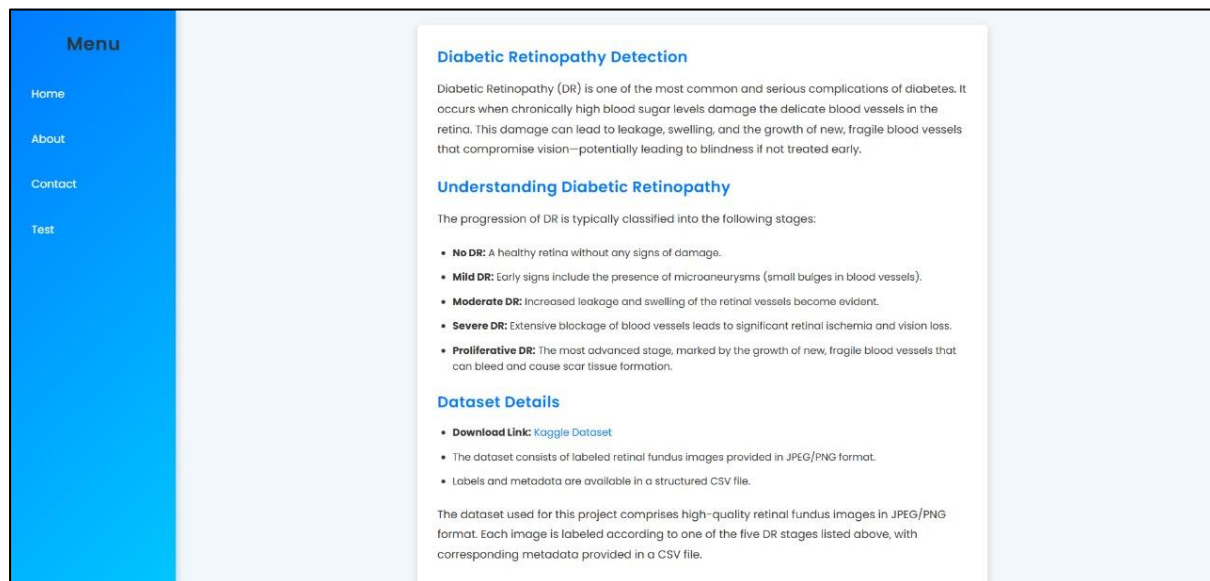
To provide a user-friendly interface for interacting with the model, we developed a frontend with four modules: **Home, About, Contact, and Test**. The **Test** module serves as the core functionality, allowing users to upload images and receive classification results.

### Key Features:

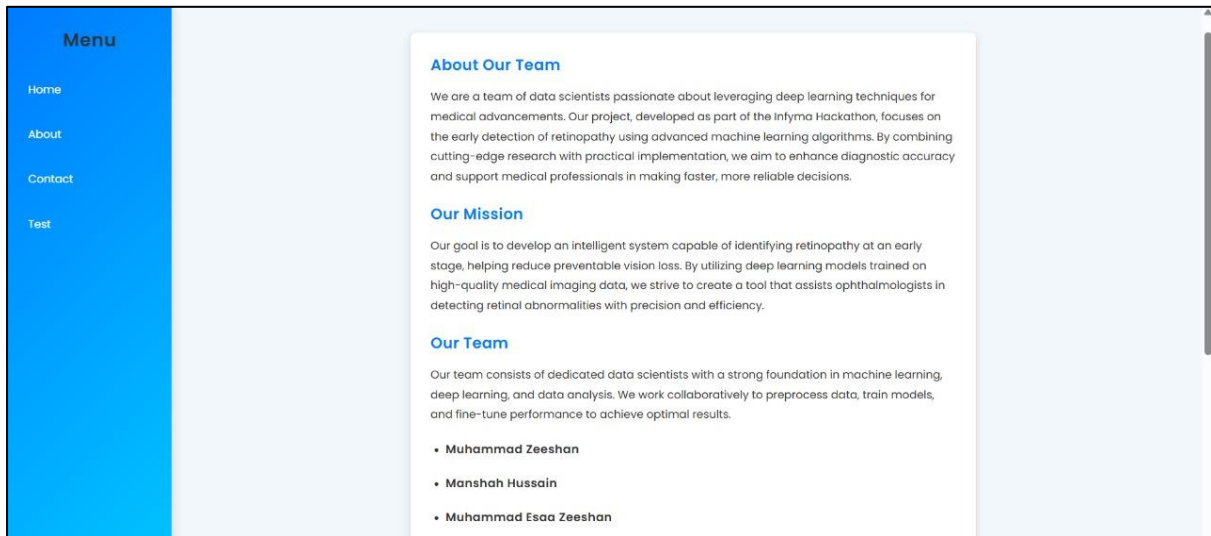
- **Home:** An overview of the application, guiding users on its purpose and functionality.
- **About:** Provides details about the project, including the model architecture and objectives.
- **Contact:** Contains information for user inquiries and support.
- **Test:** Allows users to upload images, processes them through the CNN model, and displays classification results in real-time.

The frontend was designed to be **intuitive, responsive, and visually appealing**, ensuring smooth interaction with the model.

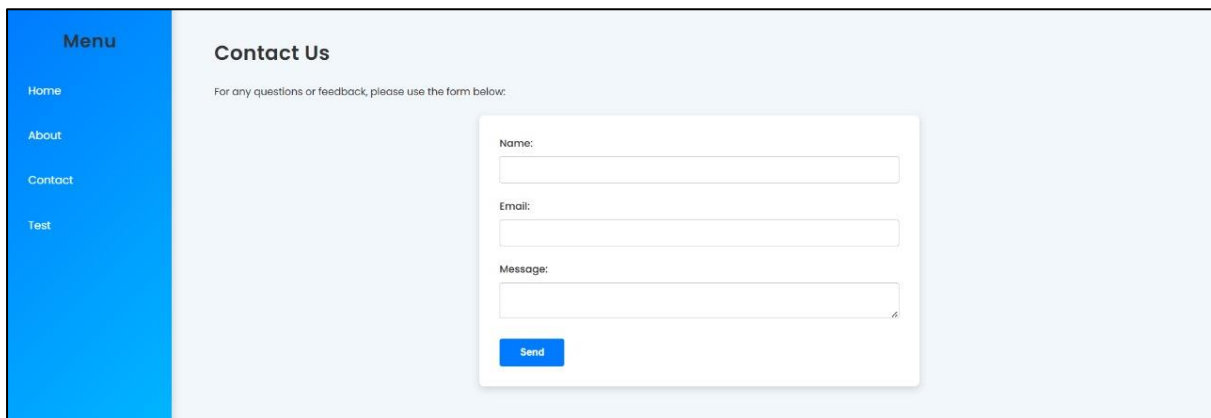
- **Home:**



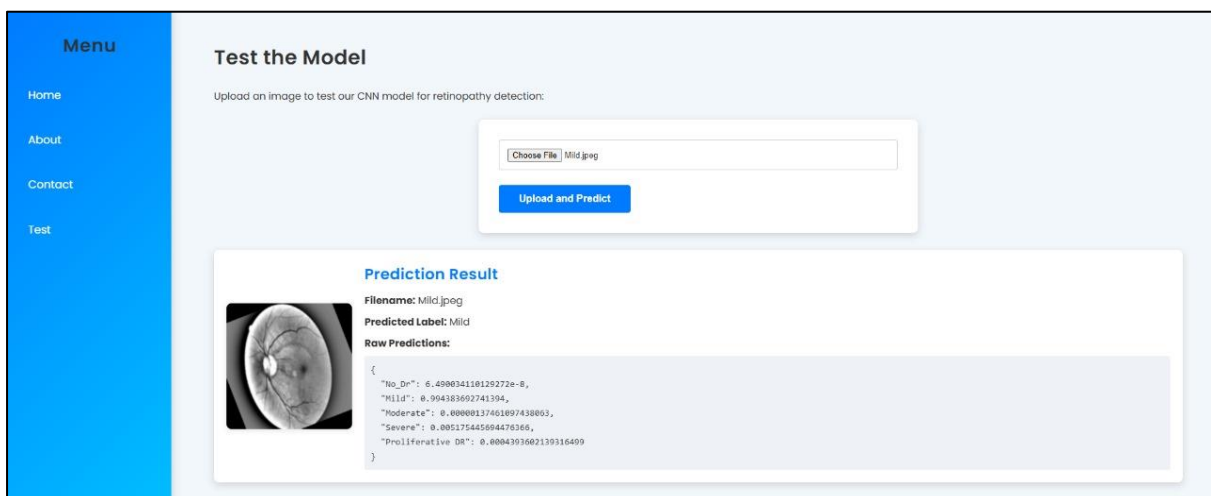
- **About:**



- **Contact:**



- **Test (Main Prediction):**



## 5. Training Process

### Compilation

The model was compiled with the following settings:

- **Optimizer:** Adam
- **Loss Function:** Categorical crossentropy
- **Metric:** Accuracy

### Model Training

The model was trained for 100+ epochs using:

- **Training Data:** 100 steps per epoch
- **Validation Data:** 20 steps per epoch

The training process outputs training and validation metrics, which are later used to assess performance.

```
base_model = ResNet152V2(include_top=False, weights='imagenet', input_shape=(224, 224, 3))

# Freeze all layers except the last 10
for layer in base_model.layers[:-10]:
    layer.trainable = False

# Define input
inputs = Input(shape=(224, 224, 3))

# Pass input through base model
x = base_model(inputs, training=False)

# Additional layers
x = Dropout(0.5)(x)
x = Flatten()(x)
x = BatchNormalization()(x)

x = Dense(256, kernel_initializer='he_uniform')(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)

x = Dense(128, kernel_initializer='he_uniform')(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)

x = Dense(32, kernel_initializer='he_uniform')(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)

# Output layer
outputs = Dense(5, activation='softmax')(x)

# Final model
model = Model(inputs, outputs)
```

```
model.summary()
```

Model: "functional\_1"

Layer (type)	Output Shape	Param #
input_layer_3 (InputLayer)	(None, 224, 224, 3)	0
resnet152v2 (Functional)	(None, 7, 7, 2048)	58,331,648
dropout_3 (Dropout)	(None, 7, 7, 2048)	0
flatten_1 (Flatten)	(None, 100352)	0
batch_normalization_4 (BatchNormalization)	(None, 100352)	401,408
dense_4 (Dense)	(None, 256)	25,690,368
batch_normalization_5 (BatchNormalization)	(None, 256)	1,024
activation_3 (Activation)	(None, 256)	0
dropout_4 (Dropout)	(None, 256)	0
dense_5 (Dense)	(None, 128)	32,896
batch_normalization_6 (BatchNormalization)	(None, 128)	512
activation_4 (Activation)	(None, 128)	0
dropout_5 (Dropout)	(None, 128)	0
dense_6 (Dense)	(None, 32)	4,128
batch_normalization_7 (BatchNormalization)	(None, 32)	128
activation_5 (Activation)	(None, 32)	0
dense_7 (Dense)	(None, 5)	165

Total params: 84,462,277 (322.20 MB)

Trainable params: 41,163,141 (157.02 MB)

Non-trainable params: 43,299,136 (165.17 MB)

### ▼ Continous Training on 5/10 Epoch Batches

```
[182]: history = model.fit(
    train_data,
    steps_per_epoch=100,
    validation_data=val_data,
    epochs=5,
    verbose=1,
    validation_steps=20
)
```

Epoch 1/5

100/100 ————— 93s 876ms/step - accuracy: 0.8421 - loss: 0.3991 - val\_accuracy: 0.7789 - val\_loss: 0.5498

Epoch 2/5

100/100 ————— 86s 862ms/step - accuracy: 0.8509 - loss: 0.3885 - val\_accuracy: 0.7742 - val\_loss: 0.5397

Epoch 3/5

100/100 ————— 87s 872ms/step - accuracy: 0.8470 - loss: 0.3925 - val\_accuracy: 0.7758 - val\_loss: 0.6074

Epoch 4/5

100/100 ————— 84s 845ms/step - accuracy: 0.8549 - loss: 0.3807 - val\_accuracy: 0.7727 - val\_loss: 0.5914

Epoch 5/5

100/100 ————— 82s 825ms/step - accuracy: 0.8523 - loss: 0.3747 - val\_accuracy: 0.7641 - val\_loss: 0.6130

```
[183]: test_loss, test_acc = model.evaluate(test_data)
print(f"Test Accuracy: {test_acc}")
```

78/78 ————— 20s 253ms/step - accuracy: 0.7009 - loss: 0.7940

Test Accuracy: 0.7724803686141968

```
[172]: # Save the model to an .h5 file
model.save('rs_model.h5')
```

### ▼ Loading Saved Model From Google Drive ¶

```
[2]: FILE_ID = "13IS0RzjePdIvhGLYoBiyw12NjQdoKU9"
MODEL_URL = f"https://drive.google.com/uc?id={FILE_ID}"
MODEL_LOCAL_PATH = "rs_model.h5"
```

```
[6]: import os
import gdown
if not os.path.exists(MODEL_LOCAL_PATH):
    print("Downloading model from Google Drive...")
    gdown.download(MODEL_URL, MODEL_LOCAL_PATH, quiet=False)
    print("Download complete.")

model = tf.keras.models.load_model(MODEL_LOCAL_PATH)
```

```
[7]: class_labels = {
    0: "No_Dr",
    1: "Mild",
    2: "Moderate",
    3: "Severe",
    4: "Proliferative DR"
}
```



## 6. Evaluation and Results

### Prediction & Metrics

After training, the model's performance was evaluated using the test dataset:

- **Predictions:**

The model generates probability distributions for each test image, with the predicted class being the one with the highest probability.


```
def predict_image(model, img_path, class_labels, img_size=(224, 224)):  
    # Load and preprocess the image  
    img = image.load_img(img_path, target_size=img_size)  
    img_array = image.img_to_array(img)  
    img_array = img_array / 255.0  
    img_array = np.expand_dims(img_array, axis=0)  
  
    # Predict  
    predictions = model.predict(img_array)[0]  
    predicted_class = np.argmax(predictions)  
    confidence = predictions[predicted_class]  
  
    # Visualization  
    plt.figure(figsize=(10, 4))  
  
    # Show image  
    plt.subplot(1, 2, 1)  
    plt.imshow(img)  
    plt.axis('off')  
    plt.title(f"Prediction: {class_labels[predicted_class]}\nConfidence: {confidence :.2f}%")  
  
    plt.tight_layout()  
    plt.show()  
  
    return class_labels[predicted_class], confidence
```

#### Testing No\_DR Image

```
img_path = "/kaggle/input/diabetic-retinopathy-balanced/content/Diabetic_Balanced_Data/test/0/10814_left_aug_16.jpeg"  
predicted_label, confidence = predict_image(model, img_path, class_labels)
```

1/1 — 0s 21ms/step

Prediction: No\_Dr  
Confidence: 0.81%



### Testing Mild Image

```
img_path = "/kaggle/input/diabetic-retinopathy-balanced/content/Diabetic_Balanced_Data/test/1/10737_right._aug_22._aug_30.jpeg"  
predicted_label, confidence = predict_image(model, img_path, class_labels)
```

1/1 ————— 0s 21ms/step

Prediction: Mild  
Confidence: 1.00%




### Testing Moderate Image

```
img_path = "/kaggle/input/diabetic-retinopathy-balanced/content/Diabetic_Balanced_Data/test/2/11156_left._aug_27.jpeg"  
predicted_label, confidence = predict_image(model, img_path, class_labels)
```

1/1 ————— 0s 42ms/step

Prediction: Moderate  
Confidence: 0.90%




### Testing Severe Image

```
img_path = "/kaggle/input/diabetic-retinopathy-balanced/content/Diabetic_Balanced_Data/test/3/11031_right._aug_5._aug_8._aug_5.jpeg"  
predicted_label, confidence = predict_image(model, img_path, class_labels)
```

1/1 ————— 0s 20ms/step

Prediction: Severe  
Confidence: 1.00%




### Testing Proliferative Image

```
[89]: img_path = "/kaggle/input/diabetic-retinopathy-balanced/content/Diabetic_Balanced_Data/test/4/1099_right._aug_22._aug_8._aug_19.jpeg"
      predicted_label, confidence = predict_image(model, img_path, class_labels)
```

1/1 — 0s 21ms/step

Prediction: Proliferative DR  
Confidence: 0.59%



- **Classification Report:**

Precision, recall, and F1-score for each class were computed to assess classification performance.

```
[178]: from sklearn.metrics import classification_report, confusion_matrix
      print(classification_report(y_true, y_pred, target_names=test_data.class_indices.keys()))
```


	precision	recall	f1-score	support
0	0.57	0.76	0.65	1000
1	0.78	0.67	0.72	971
2	0.66	0.56	0.61	1000
3	0.97	0.94	0.96	1000
4	0.99	0.97	0.98	1000
accuracy			0.78	4971
macro avg	0.80	0.78	0.78	4971
weighted avg	0.80	0.78	0.79	4971

- **Confusion Matrix:**

A confusion matrix was generated and visualized using a heatmap. This provides insights into class-specific performance and misclassification patterns.

```
[179]: import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import ConfusionMatrixDisplay

cm = confusion_matrix(y_true, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



## Performance Visualization

The training process was further analyzed through visualizations:

- **Loss Curves:**

Both training and validation loss were plotted across epochs to monitor convergence and overfitting.

- **Accuracy Curves:**


Training and validation accuracy were similarly plotted to assess model performance improvements over epochs..

```
j: plt.figure(figsize=(12, 5))

# Plot Loss
plt.subplot(1, 2, 1)
plt.plot(epochs, train_loss, label='Training Loss', marker='o')
plt.plot(epochs, val_loss, label='Validation Loss', marker='o')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training & Validation Loss')
plt.legend()
plt.grid()

# Plot Accuracy
plt.subplot(1, 2, 2)
plt.plot(epochs, train_acc, label='Training Accuracy', marker='o')
plt.plot(epochs, val_acc, label='Validation Accuracy', marker='o')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Training & Validation Accuracy')
plt.legend()
plt.grid()

plt.show()
```



## 7. Conclusion

This report outlined the development of a robust image classification model using transfer learning with ResNet152V2. Key enhancements included:

- **Data Augmentation:** Increased dataset variability to improve generalization.
- **Model Fine-Tuning:** Utilized pre-trained features while refining deeper layers for better adaptation.
- **Regularization Techniques:** Applied dropout and batch normalization to prevent overfitting.
- **Evaluation:** Assessed performance using classification metrics and visualizations.

Future improvements could focus on hyperparameter tuning, experimenting with deeper architectures, or incorporating advanced augmentation techniques to further enhance accuracy. This approach demonstrates the effectiveness of transfer learning in image classification and lays a foundation for further advancements in model performance.

---